

Witness-CL Technical Report: Constructive Contracts and Historical Evidence

Samuel Mausberg

September 2026

Abstract

A deployed agent should improve from its own experience while preserving useful earlier behavior. We study why this requires more than accumulating memory or accepting a checked update. Witness-CL combines executable learning mechanisms with explicit contracts for evidence retention, model coverage, policy continuations and exploration cost. Under supplied finite models, resets and unchanged routing, these contracts preserve incumbent value and bound cumulative exploration deficit. Online neural isolation preserves routed old functions, while representation growth requires replaying observations that were redundant in the old hypothesis class.

Controlled experiments reveal both useful learning and its limits. Reusing learned features improves reward by 12.2 percentage points over a matched no-reuse learner within a supplied grammar. Under a common query budget, however, that learner attains 47.44% final reward against 100% for simple history controls. Model-generated SQL exposes a different failure: a program can reconstruct its source answer while the reconstruction ignores the program. We implement bounded repair and a paid empty-relation intervention, and provide a Lean counterexample showing why finite reconstruction plus dependence still does not establish transfer. On a GH200, a 45-minute follow-up admits five relations but never executes the checked reuse path before its deadline. Stored guards sometimes answer fresh queries; warm qualification does not prevent stale scalar errors. Retention panels remain incomplete. The resulting case study identifies concrete computational, statistical and representation barriers, with source-linked experiments and conditional proofs. It motivates a falsifiable delayed-corroboration mechanism. General stateful-benchmark superiority and unrestricted no-forgetting remain unestablished.

1 The missing result

An online agent can retain every byte of a successful interaction and still perform worse later: retrieval may select irrelevant evidence, compression may drop a convention, or a check may consume the last tool call needed to solve the task. Conversely, a capable model can answer every question without learning anything specific to

the stream. The research target is therefore a conjunction: experience must cause useful transfer, the complete system must beat competent controls under a declared resource constraint, and protected old behavior must remain adequate.

Continual Learning Bench (CL-Bench) makes this distinction concrete with stateful-versus-stateless gain across six domains. Its evaluated full-context ICL systems are strong controls [1]. AgentCL separately shows the value of controlled compositional streams for diagnosing reuse [13]. Neither benchmark’s motivation licenses the universal claim that memory systems never help. The question is which update mechanism helps on which stream, at what cost, and with what regressions.

Witness-CL investigates a narrow candidate: learned, parameterized SQL relations whose admission is supported by the agent’s own executed observations. Its SQL fixture exposes opaque names and a queryable documentation catalog, allowing conventions to be learned across episodes. It is a custom mechanism fixture, not a native CL-Bench result. This distinction matters because a single catalog read can expose several conventions at once, and later direct SQL can already attain the one-query floor.

This consolidated paper offers three bounded contributions. First, it makes evidence retention, model coverage, continuation equality and exploration cost explicit in executable controllers and checked conditional contracts. Second, it measures representation learning, neural isolation, planning and feature reuse against controls that expose their assumptions and costs. Third, it connects model-generated program admissions to exact execution traces, repairs an observed false admission and specifies the missing transfer test. The mechanisms are related experimental systems, not one trained general-purpose agent. Useful restricted results coexist with failed stronger comparisons; general benchmark superiority and unconditional no-forgetting remain open.

2 Learning and retention

For the current SQL system, let H_t be the agent’s legal history through episode t , M_t its persistent state, and θ a fixed pretrained model. Deployment updates are

$$M_{t+1} = U_\theta(M_t, H_{t+1}), \quad \pi_t = \pi_\theta(\cdot \mid M_t). \quad (1)$$

For the current SQL agent, learning is nonparametric: program proposals, observations and memory updates occur between episodes; model weights do not change. The absence of an offline retraining phase does not make these updates free. π_t includes retrieval, prompts, decoding, routing, tool limits and fallback behavior.

For an episode instance x , let $A(\pi, x) \in \{0, 1\}$ denote correctness and $C(\pi, x)$ record SELECTs, model tokens, latency and retained bytes separately. Define accuracy on a fixed distribution D by

$$V_D(\pi) = \mathbb{E}_{x \sim D, \xi} [A(\pi, x; \xi)], \quad (2)$$

where ξ includes solver randomness. A deployment gain compares the learned state against the same backbone with its state reset, on matched instances and resource rules. A positive gain alone does not show an advantage over full history. Costs are reported jointly with accuracy; an unequal-accuracy token table cannot establish matched-accuracy efficiency.

Definition 1 (Scoped retention). *Choose a protected reference policy π_{ref} , distributions $D_1, \dots, D_K$, allowed losses $\varepsilon_k \geq 0$, and error probability α before evaluation. A distributional retention claim requires*

$$\Pr[\forall t, k : V_{D_k}(\pi_t) \geq V_{D_k}(\pi_{\text{ref}}) - \varepsilon_k] \geq 1 - \alpha. \quad (3)$$

This is stronger than a final-checkpoint result, but weaker than a pointwise claim about every realized task and random seed. Our experiments establish neither. A frozen checkpoint evaluated once can support only that checkpoint’s specified distributional comparison. Repeated promotion requires a valid lifetime error allocation; repeatedly accepting a small loss against the last policy can accumulate a large loss against the original reference.

Preserving a program’s bytes does not imply Eq. (3). Even identical outputs do not protect reward when new checks change latency or resource use. Under hidden concept drift, preserving an old answer can be wrong. One must distinguish retaining competence on a fixed protected distribution from adapting to a newly defined one.

Why a universal claim is unavailable. If two environments produce identical legal histories but require different answers on a new question, an agent with that history must issue the same answer distribution in both. It cannot guarantee correctness in both worlds. For example, a zero-valued observation is consistent with different monetary scales. This is an elementary identifiability boundary, not a claim that all structured environments are unlearnable. Restrictions on tasks, informative observations and explicit statistical assumptions are essential.

3 Antecedents and the comparison we owe

Executable reuse and iterative repair are established ideas. Voyager accumulates executable skills and uses environment feedback to improve them [18]; Agent Workflow Memory induces reusable routines, including during online interaction [20]. Dynamic Cheatsheet retains transferable text and code [15]. Thus neither a program library nor retries is our novelty claim.

ACE addresses information loss during context rewriting [23]; Evo-Memory’s ReMem integrates reasoning, action and memory refinement [22]; MemProbe combines interactions, insights and skills [13]. These motivate controls that can retain working SQL and exact documentation. Our old evolving-text arm is a particular implementation, not a faithful replication of those algorithms.

Memento 2 formalizes reflected decision processes, while Memento-Skills uses evolving external skills [19, 24]. The former’s asymptotic connection to optimal behavior assumes local consistency of the LLM near relevant cases and improving coverage; its policy-iteration and memory convergence results impose additional conditions. Those premises do not supply a finite-budget retention guarantee for this solver.

Neural memory and test-time weight adaptation provide a different axis. Titans learns a neural memory mechanism [2]; end-to-end test-time training updates on context through next-token prediction and meta-trains updateability [16]. Their sequence-modeling results are not evaluations of the present SQL policy. A frozen-backbone study cannot establish superiority over these trained architectures. A later weight-update comparison must hold available experience and total computation explicit.

High-confidence policy improvement and confidence sequences already provide routes to statistical protection under sampling assumptions [8, 17]. They do not discover useful representations, make audit data free, or predict arbitrary future drift. Our potential contribution lies in the measured connection from evidence to computation to transfer, with these boundaries exposed.

Predictive state representations already model latent dynamics through predictions of observable experiments [14]. Baseline bootstrapping constrains policy changes with limited evidence [10], and DeepSPI couples representation/world-model learning to safe online policy improvement [3]. Thus our finite-class contracts are not a first solution to representation-aware policy improvement. Their role is to expose exact coverage, continuation and resource obligations in executable form.

For the SQL component, query-equivalence tools are a stronger antecedent than instance agreement. SQLSolver proves supported equivalences using integer arithmetic, while VeriEQL checks bounded database equivalence and

can produce counterexamples [4, 7]. We have not integrated either. Their guarantees depend on supported syntax, constraints and modeled SQL semantics; equivalence to an observed source query still does not prove that the source answers the intended question.

4 Evidence and execution contracts

The earlier constructive systems answer a narrower question than the SQL learner: what can be guaranteed when the environment family or a valid state bound is supplied? Their value is to isolate the assumptions that disappear when a language model must discover the representation itself. The following contracts apply to those systems; they are not silently inherited by the current SQL agent.

4.1 Evidence must survive representation growth

For a hypothesis class $\mathcal{H}$ and exact executed observations E , define

$$\mathcal{V}(\mathcal{H}, E) = \{h \in \mathcal{H} : h \text{ agrees with every } e \in E\}. \quad (4)$$

In a fixed finite class, keep an observation as a compact witness only when it strictly shrinks the current version space. Induction gives the same final version space as retaining every constraint, with at most $|\mathcal{H}| - |\mathcal{V}|$ such witnesses. This equality does not require realizability; retaining the true hypothesis additionally requires correct feedback and its initial inclusion.

The qualification “fixed class” is essential. If new hypotheses $\mathcal{G}$ are introduced, the correct update is

$$\mathcal{V}(\mathcal{H} \cup \mathcal{G}, E) = \mathcal{V}(\mathcal{H}, E) \cup \mathcal{V}(\mathcal{G}, E). \quad (5)$$

Old constraints discarded as redundant on $\mathcal{H}$ can reject members of $\mathcal{G}$. For example, if the old class contains only the constant-zero rule, observing zero at input one removes nothing. A later rule predicting one there is nevertheless inconsistent with that observation. Replaying only the compact old witness set would admit it. Representation growth therefore stages the enlarged class, replays complete original evidence and commits atomically. An archive preserves the possibility of reinterpreting observations; it does not certify that the enlarged class contains the true rule.

4.2 Hidden dynamics with an exact partial-table cover

A stronger implemented system assumes a stationary deterministic transducer with at most K hidden states, known finite action and output alphabets $\mathcal{A}, \mathcal{O}$, and genuine episode resets. It observes executed input/output

traces, not hidden states or supplied transition tables. Public objectives g specify bounded reward functions of outputs. Policies are frozen finite trees over observed output histories.

The learner maintains a union of partial transition tables. A defined edge survives an observation only if its output agrees. An undefined visited edge branches over its K possible successor states while fixing the observed output. End states are tracked inside an episode and discarded at reset; identical partial tables are deduplicated. Every consistent complete table follows one retained branch, and every retained completion agrees with the observations. This is exact finite-class inference, not learned validity of the state bound. Selecting a minimum-size fitted automaton would lose the coverage property needed by the checker.

For candidate p , incumbent b and current consistent class C , paired symbolic execution computes

$$L_C(p, b; g) = \min_{m \in C} \{V_m(p, g) - V_m(b, g)\}. \quad (6)$$

Both executions use the same unknown-edge assignments. A branch may close when the complete remaining program, symbolic state and time agree; unread edges then affect both continuations equally. Equality of the current action or accumulated reward alone cannot justify closure. Exact residual-program keys avoid equating unequal suffixes merely through a hash collision.

Proposition 1 (Conditional incumbent preservation). *Within a stationary era, suppose the actual transducer remains in every complete class C_n , the programs and dispatch semantics are immutable, and there is a separate incumbent $b_{n,g}$ for each public objective g . Only a completed comparison with $L_{C_n}(p, b_{n,g}; g) \geq 0$ replaces that objective’s entry; all other entries remain unchanged. Then incumbent reset value is nondecreasing for each protected public objective, regardless of how an online neural or symbolic proposer changes.*

Proof. Substitute the actual transducer into each universal inequality in Eq. (6), then use transitivity across accepted replacements. Class contraction preserves an earlier universal obligation. Class expansion need not; it requires recertification and an explicitly new trust era. $\square$

The preserved Lean statements formalize the implication from coverage and bounds to retention. The full optimized Python search is not verified in Lean. A work cap returns UNKNOWN, never a certificate from a truncated class. An empty class returns inconsistency, avoiding vacuous certification.

This machinery is expensive. Along a paired path, at most $u = \min(K|\mathcal{A}|, 2H)$ previously unknown edges can be encountered. Branching can reach $(K|\mathcal{O}|)^u$ assignments per partial table, before cover and horizon overhead. A complete contingent policy tree can itself

have $\sum_{t=0}^{H-1} |\mathcal{O}|^t$ nodes. Common suffixes save work on particular programs without making inference polynomial.

When full state is observable. The earlier supplied-model controller uses the stronger all-state condition

$$r_m(t, s, c) + V_m^b(t+1, T_m(t, s, c)) \geq V_m^b(t, s) \quad (7)$$

for every live model, time and state, with $c = c(t, s)$. Backward induction then replaces the old continuation value by the no-smaller candidate continuation, yielding $V_m^c(t, s) \geq V_m^b(t, s)$ throughout the horizon. This is ordinary Bellman policy improvement, formalized in `bellman_improvement`. It is a sufficient local condition; jointly beneficial changes can fail it. The later hidden-state comparison instead checks complete policies from reset and must not be described as establishing all-state improvement.

4.3 Exploration has a separate contract

An incumbent can be uninformative. Before executing an exploratory policy p_n , the controller computes a nonrefundable debit

$$d_n = \max\left(0, \max_{m \in C_n} [V_m(b_{n,g_n}, g_n) - V_m(p_n, g_n)]\right). \quad (8)$$

It executes only if cumulative debits remain within a declared B . Under the same coverage, deterministic reset and reward assumptions,

$$\sum_{n=1}^N [V_{m^*}(b_{0,g_n}, g_n) - R_n] \leq \sum_{n=1}^N d_n \leq B \quad (9)$$

for every completed prefix. The first inequality follows episode by episode from preservation of the incumbent for the executed objective and the worst-case debit, with objective-specific reference b_{0,g_n} . Favorable outcomes do not refund the announced debit.

Equation (9) permits individual losses. It compares each execution with a reference beginning at the same reset, not with the counterfactual lifelong trajectory of a reference agent in a world with cross-episode effects. Resetting the ledger at every detected regime does not establish one lifetime bound. Unknown hidden drift can invalidate a certificate before feedback reveals the change. These distinctions separate reward accounting from general safety.

4.4 Protect the closed system, not just its module

Let $F(z, u)$ be a full system transition, where z includes controller memory, routing, normalized inputs, permissions and environment state. If old and new systems have equal $F(z, u)$ on a protected set P , and their common

successor remains in P , induction on the input tape gives equal output traces and final states. The closure condition prevents a system from leaving the protected set after one matching action and diverging later. The corresponding Lean contract is `closed_continuations_equal`.

This is the precise reason that frozen modules can be useful but insufficient. An unchanged module with a new router or input normalizer need not preserve the old computation. The SQL runtime currently supplies neither a verified closed state invariant nor the ordinary reference answer after a failed check.

5 Learning representations and parameters

The repository contains several online update mechanisms. They operate under different information contracts and must not be described as one end-to-end trained language agent. The unifying design is to let a mutable learner propose a change while a separate, explicit mechanism controls its use. That division prevents a proposal score from becoming its own certificate.

5.1 Projected residuals and their capacity cost

An early parameter-learning component freezes feature map ϕ and predictor f . If the columns of U are an orthonormal basis for protected feature vectors, set $P = I - UU^\top$ and train only W in

$$f_W(x) = f(x) + WP\phi(x). \quad (10)$$

For $\phi(x) \in \text{span}(U)$, the residual is exactly zero for every W . This is an algebraic fixed-feature result, related to orthogonal update methods; it is not a guarantee for an arbitrarily changing nonlinear network [6, 11]. Floating-point SVD and SGD give a measured drift. The real-arithmetic residual using an approximate projector is bounded by $\|W\| \|\tilde{P}\phi(x)\|$; this does not bound additional floating-point evaluation error. No complete rounding analysis or bitwise real-arithmetic equality is claimed.

Protection consumes capacity. If protected vectors span the whole feature space, $P = 0$ and no residual update can change any output. For isotropic new features and desired linear change $v^\top \phi$, the best possible residual squared error is $\|U^\top v\|^2$: the protected component cannot be fitted. This identifies when additional observable context or new isolated features are necessary. It cannot resolve a conflict hidden from the input.

5.2 Frozen learned modules and routing

A separate neural experiment trains a new feature map and head online, then freezes that module at registration.

Earlier public-scope dispatch and input construction stay fixed. This preserves old module functions exactly under that routing contract, while consuming storage for additional modules. Architectural isolation and reuse have direct antecedents in progressive networks [12]; neither freezing nor a growing archive is novel by itself.

The diagnostic uses two possible actions and exact binary success. Chosen action plus success determines the binary target, so training does use real feedback. That reduction fails for noisy success or more than two actions without further supervision. These small neural modules are distinct from the frozen language models in the SQL studies. Their historical results quantify the isolation-versus-replay tradeoff in Section 9.1.

5.3 Numerical feature search with paid promotion

The relational learner removes the enumerated world class but supplies an 84-feature grammar. Each observed SQLite join produces rows $z_i \in \mathbb{R}^6$ and features

$$\phi_m(z) = \sum_i \prod_{k=1}^6 z_{ik}^{m_k}, \quad m \in \mathbb{N}^6, \quad \|m\|_1 \leq 3. \quad (11)$$

Reports are one- or two-term integer combinations in this grammar. The learner receives a public report token, fixed table measurements and, only after its ordinary prediction, the exact scalar target. It does not receive the hidden formula. This is stronger feedback than the later SQL agent’s correctness bit.

A growing learner begins with count and six linear sums. Once data outnumber active features by two, it projects candidate columns off the current design, scores residual correlation for at most 16 candidates, adds a feature and refits scaled least squares. Each report has at most 12 features. Promoted nonlinear syntax enters a shared bank of at most 32 features and changes search order on later reports; coefficients still come from each report’s own examples. The no-reuse arm removes only bank priority. Three-step matching pursuit and full 84-feature ridge on legal history are stronger simple controls.

This experiment learns feature selection and coefficients online, not new operators. A fourth-degree parity on the complete sign cube is orthogonal to every supplied degree-at-most-three monomial. No amount of auditing repairs that representational misspecification. Irreversible early feature choices can also fill the width cap before useful terms are selected. The move to model-proposed SQL is motivated by these limits, without claiming that it has already overcome them empirically.

6 Statistical admission and its opportunity cost

When a finite correct model class is unavailable, a frozen proposal can instead be compared with an incumbent using fresh paired outcomes. For bounded rewards, let $D_t = R_t(p) - R_t(b) \in [-1, 1]$. With a predictable stake $\lambda_t \in [0, 1]$, the process

$$E_0 = 1, \quad E_t = \prod_{i=1}^t (1 + \lambda_i D_i) \quad (12)$$

is a nonnegative supermartingale under $\mathbb{E}[D_t | \mathcal{F}_{t-1}] \leq 0$. Here $\mathcal{F}_0$ includes every observation used to select the frozen comparison, and the filtration includes all preceding training, selection and audit data. The conditional expectation of each factor is at most one. Ville’s inequality bounds ever reaching $1/\alpha_j$ by α_j . Spending $\alpha_j = \alpha/[j(j+1)]$ over all started comparisons bounds the probability of any false acceptance by α , under the stated conditional nulls [8, 21]. These are standard statistical tools, not new concentration results.

For a claim about future policy value, fresh independent sampling from a stationary protected distribution is a sufficient condition. A candidate, incumbent, routing state or scope changed after seeing its audit data requires a new valid comparison. Distinct sample identifiers alone cannot establish freshness. Adaptive exact reuse of a hidden test set is also not justified by calling it a holdout [5].

The error allocation is per gate/run, not simultaneous over all experimental arms and seeds. The implemented relational gate fixes $\lambda = 1/2$, caps audits at 64 pairs, and checks floating-point log wealth against exact rational capital for the represented rewards. Every started audit spends its error allocation, even when abandoned. Audit observations cannot enter fitting; promoted feature syntax can influence later proposals, which require fresh audit data. Its journal can detect inconsistent edits, but an unkeyed hash chain cannot authenticate rewards or prevent a complete rollback without trusted external state.

6.1 Null validity does not imply useful power

The first comparison receives $\alpha_1 = .025$. A constant gain d can pass within 64 pairs only if $(1 + d/2)^{64} \geq 40$, or approximately $d \geq .118664$. A constant five-point gain cannot pass this capped test. Worse, consider independent $D = +1$ with probability .55 and $D = -1$ otherwise. The mean gain is .1, but

$$.55 \log 1.5 + .45 \log .5 \approx -.08891 < 0. \quad (13)$$

Moreover $.55\sqrt{1.5} + .45\sqrt{.5} < 1$, so $\sqrt{E_t}$ is a nonnegative supermartingale under this positive alternative. Its

eventual crossing probability is at most $\sqrt{\alpha_1} < .159$, even without a sample cap. This is a power failure of the fixed large stake, not a violation of null validity. A prespecified smaller bet, predictable variance-sensitive betting, or a valid fixed mixture is a concrete standard alternative.

Statistical cost is not merely an implementation defect. In the independent $\{-1, +1\}$ family with mean $\mu > 0$, a sequential test with null admission probability at most a and power at least $1 - b$ satisfies the standard change-of-measure requirement

$$\mathbb{E}_\mu[N] \text{kl}((1 + \mu)/2 \| 1/2) \geq \text{kl}(1 - b \| a), \quad (14)$$

for $0 < a < 1 - b < 1$ and a stopping time with finite expected duration under the alternative; the terminal decision must be measurable at stopping. It follows by the likelihood-ratio chain rule and data processing to the final decision [9]. At $a = .025$, $b = .2$ and $\mu = .05$, it requires about 1964 expected informative observations. This restricted lower bound does not cover every reward distribution; it explains why small noisy gains are harder than nearly deterministic improvements.

6.2 What Lean adds to statistical protection

For a finite population of C contexts, let $T(p)$ sum a policy’s integer returns in $[0, B_r]$. For accepted update j , define its actual violation

$$F_j = \mathbf{1}\{T(p_j) + \tau_j < T(p_{j-1})\}, \quad \tau_j \geq 0. \quad (15)$$

The Lean accounting bridge proves

$$T(p_0) \leq T(p_n) + \sum_j \tau_j + CB_r \sum_j F_j, \quad (16)$$

and the analogous inequality after any historical check-point. Divide by $C > 0$ for the uniform-population mean. Zero allowances and zero actual violations imply no loss against any historical value. F_j is defined using true population returns, not identified with an observed audit rejection. Lean proves the deterministic accounting; it does not formalize the probability space, conditional sampling law, supermartingale property or Ville’s inequality.

The limits are visible in tiny examples. Population totals $4 \rightarrow 3 \rightarrow 2$ can each pass an allowance of one while losing two overall. Return vectors $[1, 1]$ and $[0, 2]$ have the same mean but different performance on the first context. Statistical mean protection is therefore neither pointwise protection nor automatic protection after the task distribution changes.

7 System and legal evidence

Each episode supplies a scalar question and a visible physical schema. The read-only SQLite catalog explains

semantic conventions such as monetary units, NULL handling, join keys and current-profile selection. Its contents are available only after a charged query. No gold numerical target, task identifier or evaluator recipe is passed to the learner. SQL syntax and the pretrained model remain inductive biases.

The agent receives up to eight SELECT attempts, counting errors and learning checks, and ten solve-action calls per episode. Tool results have fixed row, byte and execution limits. The authority boundary rejects writes, external databases and disallowed functions or virtual modules. Model-created SQL uses the same executor as ordinary actions. This constrains tool authority; it does not certify semantic correctness or alignment.

After submitting an answer, the agent receives binary correctness. A correct ordinary trace may support a relation proposal if a complete observed query witness is available. Post-answer learning spends the remaining episode allowance and cannot revise the submitted score. Evaluation panels freeze all learning state, including evidence stores.

7.1 Typed executable memory

A relation is represented by immutable SQL text nodes and typed parameter holes. Integers, finite real values, text and NULL have separate transport types. Binding text supplies SQLite parameters rather than SQL syntax. Inner and outer queries receive disjoint parameter namespaces during composition. The stored program is capped at 4096 SQL bytes and 32 text/hole nodes; these are transport nodes, not a bound on SQLite expression complexity. The library has at most 16 entries. A digest identifies a record; it does not authenticate its meaning.

Let q be a successful own query on database d , and let $y = q(d)$ be its complete observed scalar result. The model proposes relation $f(\cdot; b)$, observed binding b , and wrapper w . The original admission criterion executed

$$w(f(d; b), d) = y. \quad (17)$$

The wrapper could read physical tables. Consequently, it could ignore f and recompute q directly. Equation (17) is agreement on one fixture, not relational equivalence over databases.

7.2 Repair and contribution checks

The repair loop fixes the original successful observation prefix and permits at most three proposal attempts per physical session. Each new attempt receives its own prior proposal and actual parser, compiler or execution feedback. Learning checks cannot become new source witnesses. A staged candidate first reconstructs the observed scalar, then spends another SELECT on the same wrapper with the relation emptied while preserving its columns:

$$w(\emptyset_f, d) \neq w(f(d; b), d). \quad (18)$$

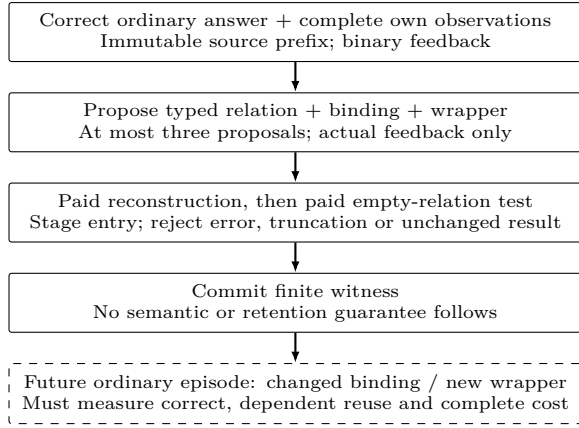


Figure 1: Evidence required at admission and at later use. The dashed box is the missing empirical milestone. Admission alone is not evidence of useful transfer.

Only complete, compatible results can establish this observed difference. An error or truncation rejects the candidate; NULL can be a valid changed result. Numerical comparisons use the existing scalar tolerance. The staged library commits only after both checks. Rejected candidates, guards and retries retain all actual costs.

An outer computation independent of the relation fails Eq. (18), fixing the observed unused-CTE admission. Useful relations can also fail this conservative test: deleting rows may leave an aggregate unchanged on a particular fixture. Passing establishes sensitivity to one intervention, not useful parameterization, correct row semantics or later transfer. The full model-driven loop has a tested implementation but no positive transfer result in the evidence reported here.

7.3 Exact-evidence control

The refined control retains exact complete observations of successful ordinary SQL executions, even if the episode’s final answer is wrong. A failed answer does not make a valid catalog response false. The final refinement separately retains every completed episode’s submitted answer and Boolean feedback, including failures. Its initial development configuration retained only successful-answer receipts; Section 9.7 preserves that distinction. The executable arm receives the same evidence store. A deterministic FIFO cap of 64 KiB bounds active evidence and program entries together; duplicate exact observations refresh their recency. Audit metadata and peak serialized retained state are reported separately. No model rewrites these observations. Learning checks and other arms’ traces are unavailable. The control therefore directly tests whether retaining observed conventions and working SQL suffices before adding program discovery.

This control has limits: FIFO can evict a rare but important observation, and exact results can be stale after drift. Its role is to remove a specific weak baseline

explanation, not to declare text memory solved. Full legal history remains a separate control, and both retain ordinary direct SQL capability.

8 Formal guarantees and limits

The current GH200 audit builds all 90 authored Lean statements, including 89 preserved statements and Proposition 2, with Lean 4.19.0. An independent imported-environment inventory checks 234 theorem constants, including generated helpers; 27 authored statements are axiom-free. Other proof dependencies are confined to the standard axioms `propext`, `Quot.sound` and `Classical.choice`. Six compiler-generated unsafe runtime specializations are separately inventoried and are not proof premises. The pinned kernel and standard library remain trusted. These counts describe verification scope, not novelty. Early manuscripts contained unchecked formal statements; repaired versions were first kernel-checked in the fifth research iteration and are checked again here. This does not retroactively verify the original archived text.

The existing typed-fragment module proves structural substitution and exact prepared-request preservation for witnessed bindings. It proves explicit CTE expansion and separation of transported values from SQL text. Giving equal requests to the same deterministic executor on the same database gives equal results by equality. This does not formalize SQLite’s parser, authorizer, optimizer or semantics. Python lexical validation and the transport boundary are checked with executable correspondence fixtures, not a complete refinement proof.

The preserved guarded-answer theorem returns a supplied ordinary answer after rejection. Its retention corollary assumes accepted protected answers already agree with the reference. Runtime fallback instead changes the prompt and reduces remaining resources; it does not receive a free ordinary answer. The conditional theorem must not be read as a runtime no-forgetting proof.

Proposition 2 (Finite dependence does not establish transfer). *For every finite witness set $W \subset \mathbb{N}$, there exists a candidate that matches reference answer 1 on every witness, changes its output when its relation is removed on every witness, and gives a wrong answer on a fresh context while still changing under that intervention.*

Proof. Let $f_W(x, 1) = 1$ if $x \in W$ and 2 otherwise; let $f_W(x, 0) = 0$. The second argument represents relation availability. Every witness gives the correct answer and intervention difference. Choose $x^* = 1 + \sum_{x \in W} x$; it is outside W , so the candidate’s nonintervened answer is $2 \neq 1$ and the intervention still gives 0. The Lean construction uses a finite list of natural numbers and proves the required fresh element and all equalities. $\square$

The result is an abstract counterexample, not a claim that the model generated this function or that arbitrary

SQL has been verified. It establishes a narrow logical boundary: any transfer guarantee needs additional restrictions on the candidate class, observations or evaluation distribution. The proof is not a sample-complexity lower bound for the implemented learner.

Resource feasibility. Let L be acquisition cost, c_h a competent direct-query cost, and $c_m = g + e$ the checked-use cost of a guard plus execution. Reusing a candidate n times saves SELECTs only if

$$L + nc_m < nc_h. \quad (19)$$

If $c_h = 1$, $g \geq 1$ and $e \geq 1$, this is impossible for $L \geq 0$ and $n > 0$. When $c_h > c_m$, the break-even requirement is $n > L/(c_h - c_m)$. This accounting observation is not a general lower bound on all memory systems: a different validator or task can have different costs. It does rule out SELECT savings for the present path against a one-query baseline. Token and latency gains remain distinct, measurable hypotheses.

9 Preserved empirical evidence

The historical tables distinguish synthetic controlled evaluations from adaptive language-model development. All language-model results below are development evidence. Adaptive diagnostics and independent follow-ups are kept separate. A replay verifies consistency with recorded SQL, prompts, transitions, source freezes and costs; it does not independently authenticate model weights or backend token counts.

9.1 Controlled learning before the language-agent study

The earlier experiments isolate capabilities that are otherwise confounded in an end-to-end agent: exact evidence reuse, recurrence recovery, neural interference, exploration choice and statistical opportunity cost. Their synthetic reward scales and feedback differ, so their scores are not pooled with SQL accuracy or normalized into a common leaderboard. Tables below report saved results; no historical model or learner was rerun for this consolidation.

Compression preserves both competence and failure. The initial fixed-class symbolic study compares witnesses with full replay over 20 seeds and 384 episodes on four public arithmetic scopes. On the interleaved scenario both achieve 94.17% whole-stream and 100% late reward; witnesses use 2340.9 mean hypothesis evaluations versus 78,146.15 for replay, with 29.4 active constraints. These are symbolic-operation savings, not measured model tokens or inference latency. Both learners also fall to 15.91% late reward under hidden shifts

and 14.09% under misspecification, producing four wrong certificates per seed in each case. Compression preserves a fixed inference rule, including its failures when coverage or stationarity is false.

Recurrence recovery. Twenty-seed arithmetic microstudies compare the original fixed-class learner, epoch resets, and resets with archive-weighted recovery. Stationary and fresh- input late accuracy reaches 100% in all three, demonstrating finite-class learning under exact feedback. Under recurring rules, archive reuse reaches 97.70% late accuracy versus 84.02% for resetting alone. Under rapid switches, those scores fall to 22.48% and 17.60%; under noisy feedback they are 65.18% and 47.62%. These stress failures violate the clean fixed-epoch assumptions needed by the exact elimination argument.

Typed relational plans show the same boundary: recurring-recipe late accuracy is 97.45% with archive reuse versus 63.59% with resets alone, but the out-of-grammar multiplication diagnostic stays below 0.11% across controls. Supported sum/difference expressions transfer to fresh rows; grammar support and semantic discovery are not interchangeable. The arithmetic and relational summaries contain both whole-stream and late accuracy; the figures here are explicitly the latter. None is a native language-agent benchmark result.

Admission opportunity cost. A separate live randomized one-step auditor, using only its chosen action’s outcome, reaches whole-stream reward .8905 on the stationary task versus .9437 for its ungated proposer, and .7544 versus .8692 under recurrence. Both stationary late scores reach one. The auditor’s useful protection and eventual learning do not repay its transient exploration and delay on these twenty-seed studies. This is an early instance of the cost problem later exposed more sharply in the SELECT-budget comparison.

Parameter learning. In 24-dimensional fixed-feature residual experiments, rank-four protection reduces new-task MSE from 25.070 to 4.226, with mean maximum protected drift 2.02×10^{-16} . Rank-twelve protection leaves MSE 12.958; full-rank protection leaves the original 25.070 unchanged. Twenty seeds per rank quantify both trainability and the capacity cost predicted by Eq. (10). This is floating-point linear residual learning, not an empirical theorem about LLM weights.

Table 1 compares learned feature modules with shared networks. The three frozen width-24 modules use 4056 parameter bytes, near the shared width-72 model’s 4040 bytes, and show zero measured old-logit drift. Replay uses an additional 10,752 replay payload bytes, excluding Python overhead, and about twice as many updates. Its late accuracy is slightly higher than isolation. All methods receive authentic public scope information; module

Neural update	Late acc.	Forgetting	Updates
Shared, width 24	96.13%	47.09 pp	4800
Shared, width 72	96.26%	44.58 pp	4800
Shared + replay	96.23%	2.65 pp	9599
Three frozen modules	95.43%	0.00 pp	4800

Table 1: Separate online neural diagnostic over 20 seeds and three public scopes. Late accuracy averages the last 200 of 1600 interactions per scope. Forgetting averages acquired-minus-later held-out accuracy across old-scope checkpoints. Fixed routing makes the frozen-module result structural. These are not language-model fine-tuning results.

storage grows as scopes are added. The 384,000 interactions are 20 seed replications across four arms, not 384,000 independent trials. Each scope has 512 fixed fresh test inputs. The historical artifact preserves per-seed summaries rather than complete interaction traces, limiting retrospective diagnosis of individual training events.

9.2 Information value and the cost of exact planning

The earlier supplied-model controller runs 17,920 episodes over 20 seeds, two finite-world domains and seven arms. With exploration budget $B = 4$, it exactly ties full-history optimistic planning in late return, averaged over the final 16 of 64 episodes, on every seed in both domains. On delayed damage, mean late return is 13.1375; requiring $B = 0$ reduces it to 12.5625. The immediate-reward greedy control reaches a worst cumulative anchor deficit of 133, compared with two for $B = 4$ and zero for $B = 0$. These are native simulator returns and deficits, not accuracy or an ICL comparison. The experiment illustrates conditional exploration protection and its cost, without establishing a reward advantage over history planning.

The finite-transducer controller with a trainable proposer improves by only .0142 reward per episode over its no-training ablation at budget $B = 64$. Informative probing improves by .0567 over unguided probing. Neither comparison establishes superiority to a strong planner: at $B = 16$, the controller trails full-history symbolic planning by .03708, with paired 95% interval $[-.05528, -.01888]$, while late return on the returning objective ties. This full-history planner is symbolic inference, not LLM ICL.

A proposed one-probe minimax-regret objective then failed its frozen test: mean reward fell .09458 relative to the earlier controller, interval $[-.16915, -.02001]$, while measured mechanics time was about $1.90\times$ larger. A post-hoc weak-dominance repair improved ten of the 256 enumerated tables without worsening one relative to that failed rule, but remained below stronger controls. The correction is a separate diagnostic, not a replacement

Finite-class selector	Reward	Plan sec.	Spent risk
Depth two	4.19688	8.3439	1.775
Informative zero-loss	3.80313	.9398	.000
Bounded optimism	4.39167	1.5603	6.975

Table 2: Forty-stream finite-class evaluation with shared work/deadline rules and $B = 16$. Equal budget ceilings do not imply equal realized risk or time. Reward is this simulator’s return per episode, not accuracy.

holdout.

The reason one-step information value can fail is concrete. Suppose two free probes reveal bits a, b and the valuable action depends on $a \oplus b$. Neither individual bit changes the immediate minimax regret, yet both together identify a better zero-loss decision. Requiring positive one-probe regret reduction stalls forever. Depth-two planning resolves this construction, but an analogous three-bit parity defeats it. The cheap rule that accepts any informative zero-loss probe reaches the correct decision after three probes. A deeper fixed horizon shifts this failure without establishing exploration completeness.

On the forty-stream comparison in Table 2, depth two beats the zero-loss rule by .39375, paired interval $[-.15845, .62905]$, but takes $8.878\times$ its planning time. It loses .19479 to bounded optimism, interval $[-.31761, -.07198]$. Only the former was primary; the other comparisons are exploratory and unadjusted for multiplicity. A cooperative search cap produces UNKNOWN on 30.73% of depth-two decisions. Smaller computation ceilings remove its primary reward gain. Exact checking remains authoritative, so these failures concern useful search and cost rather than a detected ledger violation.

A separate five-world local-LLM smoke study executes 120 model calls. The first world has no action-dependent headroom and all arms score 5.3333. On four exploratory extension worlds, the checked hybrid scores 5.5417 versus 3.9583 for raw model history and 5.0000 for a fixed-library checked control. The hybrid also has finite-model inference and an online proposer; the fixed-library control still trains that proposer. Five worlds and unequal machinery cannot establish memory superiority or model-weight learning.

9.3 Feature reuse works within a grammar, then loses its budget

The sixteen-stream relational study contains a prospectively specified positive result: prioritizing previously promoted nonlinear features improves reward on the first 24 shared-novel encounters by .12240 over the otherwise matched no-reuse grower, paired 95% interval $[-.03158, .21321]$. The first-eight contrast is exactly zero. The result supports later feature-search transfer within the supplied grammar; it does not show immediate or unconstrained discovery.

Method	Equal ordinary-example schedule			Common 4096-SELECT ceiling		
	Shared first 24	Final macro	SELECTs	Final macro	Examples	Audit pairs
Growing + reuse, audited	.1745	.6807	4889.9	.4744	138.1	582.6
Growing, no reuse, audited	.0521	.7344	4939.3	.4744	139.8	581.6
Sparse history, audited	.5130	1.0000	2440.9	1.0000	903.6	123.2
Sparse history, ungated	.5469	1.0000	1824.0	1.0000	1109.0	.0
Full ridge, ungated	.0026	.0000	1824.0	1.0000	1109.0	.0
Growing all 84, audited	.0729	.7876	4981.2	.5117	136.6	583.4

Table 3: Saved relational evaluation, sixteen independent seeds per protocol. Entries are mean arm/run outcomes. The matched schedule gives every arm 224 ordinary examples; the budget schedule exposes the audit’s opportunity cost. SELECTs exclude database setup, whose time is accounted separately.

Table 3 shows why this positive ablation is insufficient. Under the same 4096 post-setup SELECT ceiling, reuse attains .4744 final macro reward versus 1.0000 for all three fixed simple controls: audited sparse history, ungated sparse history and ungated full ridge. The paired difference is $-.52563$, interval $[-.55610, -.49517]$. Auditing consumes queries that those controls spend on ordinary learning. Ordinary execution costs three SELECTs (two measurements and one feedback query); paired audits cost five. Database creation is outside that named query budget, and wall time is not equalized.

At the end of matched training, only 34.4% of shared-novel reports contain all true terms in the selected feature set; 71.9% reach the width cap. These are evaluator-only diagnostics on saved policies, not learning hints. Old-report policies and predictions do not change during novel learning because report dispatch is authentic and frozen. That structural retention is real within its contract, but it does not test learned routing, hidden drift or shared neural parameter retention.

The full evaluation contains 192 arm/runs, 78,081 ordinary episodes, 61,790 paired audits and 61,440 panel contexts, with 727,513 post-setup SELECTs. Independent seeds, rather than those correlated episode counts, support the reported intervals. A point estimate above .05 with an interval above zero supports positive mean gain; it does not certify gain exceeding .05.

Historical evaluation boundaries. Fresh random seeds were not always new worlds. The 100-seed transducer comparison sampled 81 unique worlds, six overlapping development. The later one-probe study had 85 unique evaluation worlds and ten shared with its pilot. These are within-family evaluations, not disjoint-world generalization. A post-evaluation backend/DAG optimization preserved all selected behavior fields across 38,400 earlier rows but changed timings; timing results must retain their execution provenance. The forty-stream depth-two evaluation happened to have zero overlap with its eight development worlds, although sampling allowed replacement. Initial development preceding independently found implementation corrections remains in the archive and is excluded from the final-protocol

replay count. No corrected or post-hoc result is relabeled as an untouched holdout.

9.4 A failed prerequisite in the language-agent pilot

The v8 configuration used Qwen3-4B Q8_0 across six arms: full history, retrieved episodes, evolving insights, checked and unchecked fragments, and stateless solving. Every arm scored 0/8 on the warm block. All 288 scheduled arm-episodes completed; neither fragment arm admitted a program or performed a `USE/COMPOSE` action. Every old-task answer before and after was wrong. The 14 correct answers elsewhere in the schedule were all zero.

The pilot spent 307 SELECTs, including 141 SQL errors, and 621 model calls containing 1,138,451 reported tokens in 690.13 seconds. A lower query count for an unsuccessful arm is not a learning advantage. These results apply to that configuration, not to the feasibility of executable memory in general.

9.5 Competence improves, but memory does not yet transfer

Seven adaptive v9 diagnostics on one development seed tested interface instructions, decoding and model choice. Four 4B variants scored 0/8. Two nonthinking Qwen3.5-9B Q8_0 variants scored 1/8 and 0/8, while a thinking configuration with explicit live-tool interaction reached 7/8. Several factors changed; this is not an isolated causal estimate of reasoning or model size.

Three separately frozen follow-ups used fresh development seeds. The first stopped at native schema generation; the second was intentionally stopped after malformed memory responses and a forced reasoning cutoff. The final adapter retained host validation while repairing native schema compatibility and reasoning handling. Table 4 records its entire completed warm block.

The predeclared rule required every arm to reach at least 7/8. Text memory failed, so only 24 of 144 planned records ran. Old-task, transfer and final composition panels did not execute. Checked programs made

Method	Correct	SELECTs	Calls	Tokens
Full history	8/8	9	17	50,178
Evolving text	6/8	14	30	55,808
Checked programs	7/8	22	31	49,700

Table 4: Final v9 development warm block, seed 92003. Every arm started empty with the same solver and episode allowance. Unequal accuracies prevent a matched-accuracy efficiency interpretation. Exact evidence is a later control and has no result in this table.

seven proposals and two finite admissions, with zero later `USE/COMPOSE` actions. Its 7/8 measures ordinary solving, not executable transfer.

All seven diagnostics and three follow-ups comprise 89 records and 244 requests: 366,795 known tokens, two failures with unknown actual usage, and 1104.03 measured seconds. Declared reservations for the failed calls give an upper total of 375,904 tokens; this is a bound rather than an observed exact total. Setup and download are separate. These timings are historical and are not GH200 throughput measurements.

9.6 Two admitted programs, two different weaknesses

One admitted program computes a refunded-order scalar. Emptying it changes its wrapper result from 22 to `NULL`. However, the generator makes the count 22 by construction, and its integer parameter merely turns rows on or off. No changed-binding transfer occurred. A dependent scalar template is weaker than a reusable row-level abstraction.

The second program computes a Gold-tier total. Its wrapper repeats the physical-table query and ignores the learned relation. Emptying, removing, replacing or rebinding the relation leaves 1350.25 unchanged. The stored program itself returns the observed answer; the defect is that its admission witness does not establish its contribution. Eleven offline SQL probes on saved fixtures establish this mechanism result. They supply no learner observations and change no recorded scores or costs.

The text failures are also specific. One summary omitted the documented `NULL-as-zero` averaging convention; a later query used ordinary `AVG` and dropped `NULL` rows. Another memory state omitted units, quantities and exact join keys; the subsequent query joined incompatible identifiers and summed raw units. These observations motivate exact evidence preservation. No paired intervention isolates compression as the sole cause: the solver could have reread the catalog.

9.7 Status of the current refinement

The GH200 runner compares full history, exact evidence and repaired programs, each starting empty. Eight warm questions precede an eight-item frozen old panel, 16 further ordinary questions, the same old panel, and eight final questions: 48 episodes per arm, 144 planned records per study. Every arm must score at least 7/8 warm before proceeding; frozen panels cannot update memory. These are new development studies, not native benchmark evaluations.

The runtime uses official Qwen3-32B Q8_0 weights, a pinned CUDA llama.cpp build and one GH200 with approximately 96 GiB of device memory. The 65,536-token context uses explicit YaRN extension beyond the native 32,768 positions. Source, model and runtime inventories are bound before each study. Solving enables thinking with temperature .6, top- p .95, top- k 20, minimum probability zero, presence penalty 1.5 and sampling seed 42. Solve and nonthinking proposal calls allow at most 2048 and 4096 generated tokens. Setup and a generic transport smoke check are separate from study cost. Changed model and runtime prevent a causal cross-version attribution to memory alone.

Initial development. Stream 94000 completes 24 warm records and fails the all-arm gate: full history scores 7/8, exact evidence 4/8 and programs 4/8. All initially use integer division for currency despite the catalog’s real-arithmetic convention. History later corrects it. The evidence shelf retains successful SQL executions but omits failed-answer receipts, withholding a useful failure signal. This is a weakness of that configuration, not a general failure of exact memory. Two programs pass reconstruction and dependence; neither is subsequently executed. The run uses 61 calls, 167,024 tokens and 577.38 study seconds, with no unknown usage. Transfer and retention remain untested after the gate stops it.

Final adaptive follow-up. A separate freeze precedes seed 94001. The shelf now retains every completed answer and Boolean feedback, and every arm receives the same generic SQLite integer-division instruction. No hidden target or domain query is supplied. Both changes are selected after the first failure, so this is not a causal ablation or heldout evidence. All warm scores meet the threshold and the runner proceeds.¹ The fresh old-panel scores are 6/8, 2/8 and 4/8 for history, evidence and programs. Table 6 keeps the later comparison to common finished questions. Figure 2 shows qualification, not a transfer curve.

The follow-up stops at its deadline with 66 saved records out of 144 planned, including one interrupted program-arm episode; 65 records complete. It uses 174

¹The final manifest’s `warm_qualified` flag remains false because it additionally requires a completed study; the warm-score threshold itself passed.

Stream	Method	Episodes	Warm	All SELECTs	All calls	Known tokens	Admitted
94000	Full history	8	7/8	13	21	78,033	0
94000	Exact evidence	8	4/8	9	17	39,567	0
94000	Checked programs	8	4/8	17	23	49,424	2
94001	Full history	22	8/8	34	56	342,897	0
94001	Exact evidence	22	7/8	25	48	190,543	0
94001	Checked programs	22	7/8	51	70	247,317	5

Costs cover all recorded phases; the Warm column covers only the first eight ordinary questions. Known tokens exclude unmeasured usage.

Table 5: GH200 development receipts. These adaptive studies are separate configurations. Counts include failures, checks and the final interrupted follow-up episode; only 65 of its 66 saved records complete. The program arm’s follow-up tokens exclude one call with unknown usage. Lower costs at unequal accuracy do not establish matched-accuracy efficiency.




Figure 2: Cumulative accuracy over the eight warm questions in each adaptive development stream. The horizontal line marks the final qualification threshold, not a requirement at every prefix. These curves establish neither stateful gain against reset memory nor transfer to fresh instances.

Method	Warm	Old before	Later shared
Full history	8/8	6/8	4/5
Exact evidence	7/8	2/8	4/5
Checked programs	7/8	4/8	3/5

Table 6: Follow-up accuracy on completed matched questions. Later shared is ordinary indices 8–12, only five of the planned sixteen later questions. Index 13 is correct for history and evidence but interrupted for programs, so is excluded here and included in total costs. Failed completed attempts count as zero. Old-after and final panels have no observations.

generation attempts and 780,757 known tokens. The interrupted call has unknown usage, so the token total is incomplete and the receipt does not mark usage verified. Recorded time is 2700.024 seconds: incremental export exceeds the 2700-second ceiling by .024 seconds, and the receipt marks that cap unsatisfied. Neither later old-panel nor final-panel evaluation occurs. Five relations are admitted, but none executes through the checked `USE/COMPOSE` path on a future question. The intended

execution mechanism is not demonstrated, and retention is unassessed; missing panels are not zero-regression evidence.

Across both studies, resource expenditure is 235 calls, 947,781 known tokens plus one unmeasured call, and 3277.40 study seconds. Accuracy is not pooled across these adaptively changed configurations. The aggregate allowance is 3600 seconds, 4M tokens and 1800 attempts; exact compliance with the token allowance is not asserted with missing usage. No further model-task attempt is made in this phase, and reserved heldout seeds remain unused. Source hashes are unchanged. Source-bound replay verifies 65 complete follow-up traces and 110 recorded SQL executions, separately identifying the partial trace. This is a partial-study audit, not completion of the experimental protocol.

What the repair actually changes. On the follow-up total-refunds task, the first wrapper repeats a physical-table subquery and aliases it as `reused`, shadowing the learned CTE. Reconstruction returns 2311.75, but emptying the learned relation still returns 2311.75, so admission is rejected. The next proposal reads the

actual **reused** relation: reconstruction remains 2311.75 and the intervention returns zero. It is admitted after four charged check SELECTs plus the ordinary query. This is model-driven defect repair on one candidate, not later transfer.

Applicability is a distinct failure. The North-customer relation stores a guard requiring the old customer count, six. A later proposed composition observes three; a frozen old-panel count question observes four. Both guards reject before the relation executes. The old-panel episode nevertheless answers correctly with the guard’s fresh count as its only current SQL result. This records useful source-query reuse, without a counterfactual establishing the answer’s causal dependence on that result. Zero proposed-relation executions therefore does not mean zero memory influence. The attempted composition also contains an unsupported substitution placeholder; its validity is unestablished.

Finally, the old panel exposes stale scalar replay. History, evidence and programs answer without SQL on two, six and four questions respectively; zero, one and one of those answers are correct. The interface changes rows without exposing a snapshot identifier, while reusing question text and schema. This makes freshness an interface as well as a memory issue. Exact retained results and a successful warm gate do not ensure fresh-instance competence, and no paired intervention here isolates one cause of these failures.

10 Falsifiable next mechanisms

10.1 Ground the answer in a current execution

The interface should first make observation freshness explicit: the fixture can change rows while reusing a question and schema, without exposing a database snapshot identifier. A cheap prospective control supplies that identifier and a generic statement that earlier scalar results may be stale. This changes no hidden target or convention. It must be compared separately with the stronger intervention below, because stale answers need not be a memory-only failure.

A separate proposed repair makes the answer action refer to a complete numeric result from a query executed in the current episode. The host resolves the receipt and emits the stored value; the model cannot supply a replacement number or a receipt from a previous episode. This would apply equally to every arm, including full history, and its query costs remain charged. It targets stale scalar replay rather than reusable representation learning.

The operational hypothesis is that current execution receipts reduce wrong answers copied from previous fixtures at useful total cost. Test the interface change on independently selected development streams, compare

with an identical free-number answer interface, and preserve every error and added query. Kill its practical claim if accuracy does not improve enough to repay those costs. A constant query can still reproduce an old scalar; requiring a physical-table read would close only the simplest loophole. Neither a read nor a current receipt establishes that the query answers the question. This is an unimplemented experiment, not a theorem of semantic correctness.

10.2 Corroborate a relation before promoting reuse

Hypothesis. A useful executable memory should preserve a relation that supports multiple future operations, while text retains the exact evidence needed to interpret it. We propose keeping a relation provisional until an independently encountered future episode corroborates it with a meaningful changed binding. This is an unimplemented extension beyond the bounded repair, not an empirical result.

The learner first proposes a row-level relation from a correct own query, retaining its output columns and source computation. Admission continues to require charged reconstruction and dependence. A restricted outer language would read only the supplied learned relation, literals and approved pure operations; physical table access remains in the inner program. Enforcing that restriction requires a real SQL semantic-access boundary, not substring checks. It prevents the specific independent recomputation, but cannot certify intent.

A useful admission ablation would apply a supported-fragment equivalence checker to the complete source and reconstruction, using no hidden target query. Bounded counterexamples should be replayed in SQLite to detect dialect or semantic-model mismatches. Record proof, counterexample and unknown separately; an unsupported feature or timeout is not an admission. A CPU solver budget is charged alongside model cost. Kill this addition if supported coverage is too small, counterexamples fail to reproduce, or checking cost erases useful later gains. Even an unbounded source-equivalence proof would protect only that source computation, leaving intended semantics and new operations unproved.

Applicability needs its own repair. Current guards can require an old aggregate to remain exactly unchanged. That can reject a valid relation after rows change, while an unchanged aggregate can conceal a convention change. A prospective ablation would instead compare exact observed schema and catalog conventions, leaving variable data out of the compatibility key. This uses learner-visible metadata, not evaluator scope identifiers, and its reads are charged. It still requires truthful documentation and cannot validate a bad inferred relation.

A provisional relation is not counted as useful learning. Freeze its code and source wrapper. When a later ordinary question naturally makes the same source operation

relevant, the solver first answers without using this candidate. After correct feedback, the agent may spend its remaining allowance on reconstruction with a proposed new binding, comparing it with the new own successful query. These observed answers must agree, and both the binding and observed nondegenerate output must differ from the original witness. The learner receives only its ordinary observations and binary feedback; the evaluator does not schedule a hidden task identifier as a retrieval hint. A separate later question must use a new outer operation on fresh rows and depend on the stored relation. The evaluator audits provenance, actual parameter changes, output variation and correctness after the learner has finished.

This design separates three milestones: a checked proposal, a corroborated source computation, and useful new composition. It can still fail. Two samples need not identify the intended relation; a misleading proxy can pass; a valid relation may never recur; and extra corroboration may erase every future cost saving. The outer restriction can exclude useful joins. These are reasons to run the following experiment before adding a more elaborate architecture.

10.3 Experiment and decision rules

Freeze model revision, effective decoding, prompts, all update rules, retrieval/cap policies, generator, seeds and resource ceilings before each new study. Reserve separate development and confirmatory streams, including fresh numeric data and semantic conventions where the generator supports them. Never pool the earlier adaptive diagnostics as independent replications.

The current three-arm development comparison uses full history, exact evidence and repaired executable memory. All start empty. Every arm must first pass the same competence gate. The mechanism extension then requires a separate freeze and at least the following comparisons: no persistent memory, full history, exact evidence with SQL, bounded repair, and delayed corroboration. A faithful published evolving-text method is needed before claiming an advantage over modern memory systems. An ablation giving the same extra direct queries to the control distinguishes representation benefits from additional evidence.

Use shared-reuse, nonreuse and near-match/drift streams. Measure old panels on fresh instances before and after learning with updates disabled; evaluate changed binding and new outer operation separately. Complete schedules are the unit of replication. Report each stream’s accuracy, stateful gain, query cost, input/output tokens, latency and retained bytes, including rejected proposals and checks. Do not analyze correlated episodes as independent streams.

Chosen useful effect, not an established threshold. A reasonable next target is at least five percentage points

of future accuracy improvement over the best qualified compact control, no more than two points of loss against the protected reference on old-scope accuracy, and a 20% token reduction against full history at qualified accuracy. These are separate endpoints: passing one does not compensate for failing another. They require prospective finalization after development establishes achievable accuracy and variance, before confirmatory evidence is seen. SELECT savings should be abandoned on cases subject to Eq. (19), rather than hidden inside another metric.

Kill the mechanism as a practical improvement if exact evidence matches its accuracy at lower total cost, if apparent transfer disappears after source or binding interventions, or if it rarely produces a later correct composition through the intended execution route. Zero checked relation executions would not show that visible program text or stored guard queries had no effect on the solver. Reject a retention claim whenever its predeclared noninferiority test fails; an underpowered interval is inconclusive, not proof of preservation. Any future-label leakage, uncharged repair or mutable frozen panel invalidates the comparison even if its score is high. A negative result should terminate the corresponding claim rather than trigger repeated evaluation on the same seeds.

10.4 Resources and statistical feasibility

The current GH200 supports the selected 32B inference runtime, alongside CPU SQLite execution and local artifact storage. A query-equivalence ablation would add a CPU SMT solver; online weight adaptation would additionally need optimizer state and update compute. These are distinct resource requirements.

A three-stream, three-arm, 48-episode-per-arm pilot has 432 planned records. Warm-question cost is an unreliable estimate of its full cost: later questions in the current study require several minutes of reasoning and repeated failed queries. There is no validated complete-schedule runtime estimate. A bounded future pilot could allocate at most one GH200-hour, one million tokens and 600 generation attempts per arm and stream: nine accelerator-hours, nine million tokens and 5400 attempts in total, with setup separate. These are proposed ceilings, not a promise of schedule completion. Additional causal ablations require additional arms and budgets. Reserve resources for frozen evaluation before allocating discovery work; missing required panels leave retention unassessed. Development must establish completion and useful transfer within the ceilings before a larger confirmatory commitment.

For a broader comparison, determine independent-stream count from development variance and the fixed endpoints. As a planning approximation for a paired mean difference, $n \approx (z_{1-\alpha} + z_{1-\beta})^2 \sigma_D^2 / \delta^2$; account for all endpoints and use conservative variance estimates.

For example, with $\sigma_D = 0.10$, 80% power and a one-sided 0.005 endpoint level, detecting $\delta = 0.05$ needs about 47 streams, whereas resolving a 0.02 retention margin needs about 292. These normal approximations are not guarantees, and clustered or nonstationary data require a justified model. A short development run cannot support a two-point no-forgetting headline.

For perspective, even observing no regression in 32 independent Bernoulli trials only bounds the event probability by about 8.9% with a one-sided 95% exact interval. Episodes sharing a stream do not automatically supply 32 such trials. Retention is a substantial experimental requirement, not an adjective attached to an immutable library.

11 Training, inference and deployment boundaries

The present SQL language agent changes external state during deployment. Its training signal is its own permitted observations, execution errors and binary task feedback. It performs no offline model retraining and no online gradient updates. This isolates the representation experiment, but it may leave the solver unable to discover useful relations even with a sound admission rule. Increasing model capacity can address competence; it cannot retrospectively turn ordinary SQL into learned reuse.

A later online LLM adapter is a distinct proposal: train only on independently verified own successful traces, keep an immutable reference model, and test candidate updates on fresh protected tasks before promotion. It would need optimizer state, replay/storage accounting, update compute, and a full-history control using the same examples. No such adapter or retention result is claimed here. Offline meta-training for updateability also changes the resource and supervision comparison.

During inference, rejected guards consume budget and can influence subsequent reasoning. A deployment requiring a hard service fallback must reserve enough resources for the reference path and specify its inputs and environment state. Merely retaining a reference checkpoint supplies none of those conditions. Observations and generated memory are data, not authority to modify the executor. Database read-only enforcement and deterministic source provenance reduce specific implementation risks; safe decision-making beyond those boundaries remains a separate task.

12 Limitations and research value

The strongest general comparison remains negative. The current fixture has a small set of conventions, a revealing catalog and one structurally constant target. The v8 solver was inadequate, and the competent v9 follow-up

stopped before transfer. The v9 interventions diagnose two saved programs, not a population of learned abstractions. New GH200 traces separately demonstrate repaired admissions, stale answers and guard failures. The exact-evidence control and repair address observed defects; they do not establish a general solution. The final GH200 study exceeds its wall ceiling by .024 seconds during export, has one call with unknown token usage, and lacks its later retention and final panels. None of these gaps is silently filled by a model or estimate.

The potentially useful result is a more demanding unit of evidence. A memory entry should not count as acquired competence until its computation contributes to a correct later task under meaningful variation and complete resource accounting. Even then, retention and benchmark superiority require separate measurement. A strong future paper must demonstrate that conjunction against competent full history and modern memory controls on independent stateful streams. The next decisive evidence is correct dependent reuse on completed controlled streams, followed by a qualified retention comparison.

A Recurrence, noise and limited storage

A.1 A recurrence-sensitive mistake bound

Within one deterministic, realizable, unchanging inference epoch, assign each rule h positive prior mass $\pi(h)$. Predict a maximum-mass action and delete rules contradicted by the exact chosen-action feedback. If at most $K \geq 2$ distinct actions are predicted, a mistake removes at least $1/K$ of surviving mass. The true rule keeps its original mass, so after M mistakes

$$\pi(h^*) \leq (1 - 1/K)^M, \quad M \leq \frac{\log(1/\pi(h^*))}{-\log(1 - 1/K)}. \quad (20)$$

This is the standard weighted-elimination argument specialized to recurrence. If an archive $\mathcal{A} \subseteq \mathcal{H}$ receives prior mass β ,

$$\pi(h) = \frac{1 - \beta}{|\mathcal{H}|} + \frac{\beta}{|\mathcal{A}|} \mathbf{1}\{h \in \mathcal{A}\}, \quad (21)$$

a returning archived truth has $\pi(h^*) \geq \beta/|\mathcal{A}|$ and therefore a smaller recovery bound than a uniform new search when the archive is useful. An empty archive uses the uniform prior. The reference uses $\beta = 4/5$.

Contradiction first triggers bounded grammar growth with complete evidence backfill; if no consistent interpretation remains, an internal epoch restarts while old programs stay archived. The epoch is not a supplied hidden regime identifier. A uniquely identified rule in a fixed class encounters a contradiction on its first wrong post-change prediction, but the argument does not cover

a still-ambiguous class, noise, rapid switches or a newly constructed mixed-history explanation. The stress failures reported in the main text are therefore substantive limits, not violations of the clean-epoch theorem.

A.2 Noise needs probabilistic evidence, not harder deletion

A separate component assumes a fixed finite collection of proper conditional outcome models that includes the true model. Fix weights $\pi_m > 0$ with $\sum_m \pi_m = 1$ before observing data, let $L_{m,n}$ be model m 's likelihood of the actual chosen-action outcomes, and define

$$E_{m,n} = \frac{\sum_j \pi_j L_{j,n}}{L_{m,n}}. \quad (22)$$

On paths with positive true-model likelihood, likelihood ratios are nonnegative supermartingales under that model, and martingales with appropriate common support. Predictable adaptive actions do not break the conditional likelihood argument. Consequently, excluding m only after $E_{m,n} \geq 1/\delta$ retains the actual fixed model at all times with probability at least $1 - \delta$.

The implementation uses exact rational arithmetic in small examples and checks all binary histories of length six for a common-support example. In 1000 independent 64-step diagnostics with 10% noise and nominal $\delta = .05$, the probabilistic filter ever excludes truth in 12 sequences, versus 997 for hard deletion. This finite simulation illustrates the distinction; it is not a proof of the nominal coverage probability. This is a separate tested statistical component, not a stochastic extension of the integrated deterministic controller or a Lean proof of probability theory. Uncalibrated neural scores are not specified conditional probabilities. Models chosen after seeing old data cannot receive retroactive likelihood credit, and stochastic realized returns need their own deviation argument beyond the pathwise deterministic debit bound.

Soft fixed-archive routing is also implemented separately through importance-weighted feedback and fixed-share expert updates. It keeps alternatives alive when hard deletion is inappropriate, but provides an expected tracking bound rather than per-episode protection. For rapid switches, the stated regret bound can exceed the whole horizon. It does not prove learned routing over an arbitrarily growing archive, and is not integrated into the language-agent run.

A.3 A storage boundary

After N independent binary facts, exact answers to every possible old-fact query must distinguish 2^N histories. A deterministic memory with fewer than N bits has fewer states, so two different histories collide and at least one old answer must fail. This pigeonhole argument

concerns unrestricted exact recall, not compressible task structure. A bounded active cache can coexist with an archive; bounded total memory cannot retain unlimited independent information exactly. The evidence FIFO cap and isolated module archive make opposite resource choices and neither alone solves this boundary.

B Local auditing and shared inference

B.1 Locality changes the cost of collecting evidence

For frozen one-step policies b, c and an observable gate g , define $p(x) = c(x)$ inside the gate and $p(x) = b(x)$ outside it. Under a valid common-input comparison, reward difference D is exactly zero outside the gate. On a stationary distribution with gate mass $\rho > 0$,

$$\mathbb{E}[D] = \rho \mathbb{E}[D \mid g = 1]. \quad (23)$$

Positive conditional gain implies positive global mean gain, while protected subgroups inside the gate still need separate obligations. A disagreement gate can itself require two costly model predictions; locality is useful only when screening is cheap or independently justified.

For a fixed-stake betting process, deleting exact-zero differences leaves the product unchanged. At N deployment opportunities and n_g gate events, dense paired execution uses $2N$ rollouts and screening uses $N + n_g$. The reduction in extra rollouts can approach $1/\rho$; the reduction in total rollouts is at most two. Neither produces additional statistical information. For predictable stakes that depend on skipped observations, equality additionally requires matching the stake process rather than silently changing its history.

In 4000 simulated replications at 2048 opportunities, a diffuse .05 mean gain passes the reference audit in 27% of runs. A sparse process with the same global gain, gate probability .1 and conditional gain .5 passes in 100%. In the sparse case, screening reduces mean extra rollouts from 344.307 to 34.45625 and total rollouts from 688.614 to 378.76325. That is about $1.82 \times$ total savings, not tenfold. These favorable low-variance comparisons motivate locality but do not imply every useful deployed improvement is localizable.

B.2 Randomization when two rollouts are unavailable

If only one real action can execute, randomize treatment $A \in \{0, 1\}$ before observing its reward. Conditional on history and current context, let $A \sim \text{Bernoulli}(e)$ independently of potential outcomes, with logged $e \in [p_{\min}, 1 - p_{\min}]$ and $0 < p_{\min} \leq 1/2$. The observed reward

Rows	Direct Python, ms	Shared cube, ms	Ratio
32	.479895	.129626	3.70
512	7.623821	.611038	12.48
8192	113.284889	6.709610	16.88

Table 7: Historical CPU primitive medians for 64 supplied plans. These are neither optimized-database comparisons nor end-to-end or GPU speedups.

is the outcome for the assigned action. For $R \in [0, 1]$, the centered estimator

$$Z = \left(\frac{A}{e} - \frac{1 - A}{1 - e} \right) (R - 1/2) \quad (24)$$

has conditional expectation equal to candidate-minus-baseline expected reward and magnitude at most $1/(2p_{\min})$. The treatment and control centering terms cancel in expectation. Using one fixed normalization preserves the intended estimand; context-dependent normalization can change it.

This estimator assumes randomized conditional assignment and meaningful potential outcomes at the same starting context. Cross-episode action effects cannot be silently ignored. The treatment action can lose reward while a candidate is still being audited. This live randomized branch therefore addresses missing counterfactual feedback, not risk-free experimentation.

B.3 Factor shared relational work before optimizing kernels

The early typed relational grammar uses F nullable Boolean fields and C numeric fields. Partition rows by their ternary patterns in $\{\text{NULL}, 0, 1\}^F$, and cache each numeric sum within a bucket. A filter selects a known subset of patterns; its sum and difference recipes follow by addition over those buckets. Disjointness and distributivity establish equality to direct interpretation, with explicit NULL and empty-table semantics checked against SQLite.

For H_r recipes over R rows, direct interpretation costs $O(H_r R(F+1))$. A cube with P_f distinct filter signatures costs

$$O(R(F+C) + P_f 3^F (F+C) + H_r) \quad (25)$$

and uses $O(3^F C + P_f C + H_r)$ storage including cached filters and materialized outputs. The exponential term makes small F a substantive condition. Checked integer bounds are needed before narrowing a SQL sum to a smaller machine type.

Table 7 shows a real factorization benefit against direct Python. The separate cold packed-continuation diagnostic does not: at 64 models it takes 10.024 ms versus 9.830 ms for the scalar reference. Shared-prefix simulation matches independent execution on 200 cases and reduces mean environment calls from 64 to 39.405,

a	b	Anchor	Either probe	Even	Odd
0	0	5	5	10	0
0	1	5	5	0	10
1	0	5	5	0	10
1	1	5	5	10	0

Table 8: Two free observations have complementary decision value. The probe reward is constant but its observed bit differs across worlds.

while controller calls remain 64. An implementation can save symbolic work without reducing model generation.

The custom evidence kernels are separate from llama.cpp’s CUDA inference. A new GH200 check compiles the filter and consensus kernels with CUDA 12.8 and passes 336 seeded fixtures, including empty sets, integer extremes, warp boundaries and partial tails. Each fixture checks both operations. The same fixtures are repeated under Compute Sanitizer memcheck and racecheck, with zero reported errors or hazards. The receipt `artifacts/v10/cuda-validation.json` binds source and binary hashes. This establishes finite kernel checks, not a GPU speedup, complete kernel proof or model-learning result.

C Complementary information defeats fixed-depth value tests

Consider four one-state worlds indexed by $(a, b) \in \{0, 1\}^2$. The anchor and two bit probes give reward five; each probe also reveals its corresponding bit. Choosing the correct parity gives ten and choosing the wrong parity gives zero. All actions and this restricted world class are available to the finite checker.

Initially the anchor’s worst-case regret is five. Either individual bit leaves both parities possible, so the best robust regret remains five after every single-probe outcome. The parity decisions each need worst-case debit five and are infeasible at budget zero. A rule requiring strictly positive one-probe regret reduction selects the uninformative anchor; its class never changes and it can remain there forever. Executing both zero-loss probes instead identifies parity and admits a gain of five.

The depth-two successor and the cheap informative-zero-loss control both achieve returns 5, 5, 10 in all four worlds. Three-bit parity exposes their difference: any two bit observations still leave parity unknown, while the cheap rule eventually takes all three. More generally, requiring positive terminal value within fixed lookahead k can miss a $(k+1)$ -observation parity structure. This is a constructive mechanism argument, not a lower bound for all exploration algorithms or a claim that exact deep search scales.

D A route to learned approximate state

An alternative to expanding exact tables is to learn an abstract predictive state and a reward/transition model. This remains a proposed extension. First assume a fixed valid state and policy correspondence: true and learned kernels must share a measurable state space, or a specified abstraction under which the true dynamics are Markov, with a common initial state. Learning this sufficiency is itself unsolved here; hidden state aliasing is not covered by the following calculation. Suppose rewards lie in $[0, R_{\max}]$ and the learned model has uniform reward error at most ϵ_r and transition total-variation error at most ϵ_p over all relevant state-action pairs. For a fixed policy and remaining horizon h , a standard simulation-bound argument yields

$$|V_{\text{true}} - V_{\text{model}}| \leq h\epsilon_r + R_{\max}\epsilon_ph(h-1)/2 =: E_h. \quad (26)$$

At each step add reward error, the next-state value error, and at most $\epsilon_p(h-1)R_{\max}$ from the transition distribution. Summing proves the bound. A model-predicted candidate advantage greater than $2E_h$ then suffices for a positive true advantage. Adaptively selecting a policy requires the error envelope to hold uniformly over that selection class.

Small training loss supplies no such envelope. A concrete learner could retain raw distinguishing traces, split an abstract state when the same encoded state predicts incompatible continuations, freeze each abstraction version and recertify affected programs. This is related to predictive-state and safe world-model learning, not a first approach of that kind [3, 14]. The representation proposal is killed if required coverage cannot support any useful update before drift, if its state count grows to raw-history size, or if a simpler history model matches its return with lower total compute.

The SQL delayed-corroboration proposal tests the analogous smaller question: can a reusable relation replace repeated reconstruction of a task’s latent semantics? It has the advantage that execution can expose wrong computations cheaply. It has the disadvantage that finite witnesses, even intervention-sensitive ones, do not supply the uniform envelope above. The constructive positive theory and the open discovery problem meet at that missing premise.

E Artifact map and reproduction

The repository keeps source, formal contracts, tests, protocols and machine-readable traces separate. The frozen v8/v9 evidence is cited by its version to avoid pooling different source configurations. The key entry points are `artifacts/v8/development-replay.json`, `artifacts/v9/results.json`, and `artifacts/v9/mec`

`hanism-counterfactuals.json`. `docs/v8/FORMAL.md` explains the compiler proof boundary; `docs/v9/RESU LTS.md` and `docs/v9/MECHANISM_REVIEW.md` give the historical trace-level diagnosis. The GH200 records are in `artifacts/v10/development-94000` and `artifacts/v10/development-94001`; `artifacts/v10/results.json` contains completed-study aggregates. The two pre-run freeze receipts bind their separate source and protocol inventories. Section 9.7 distinguishes their status. No native benchmark score is inferred from an interface compatibility test.

Saved-data checks require the repository’s Python test and analysis extras. The historical audit command is `make legacy-audit`; it invokes no model. The typed Lean build is pinned to Lean 4.19.0. The new finite-witness result is in `formal/WitnessCL/Intervention.lean`; `artifacts/v10/formal-audit.json` records the build, dependency inventory and source hashes. The current command is `python3 formal/audit_v10.py -output artifacts/v10`. Its fixtures match the preserved reference files byte for byte. This distinguishes build verification from Python/SQLite correspondence. The stable build commands are `make formal` and `make paper`. The latter uses `pdflatex`, `bibtex` and two further LaTeX passes. Historical manuscripts are retained in Git at `88b0a1b7c093c21135489a3e5cbbda3bb0634536`; the consolidated source is `paper/main.tex`. `docs/v10/HISTORICAL_EVIDENCE.md` maps every early table to its saved artifact and records overlap/correction boundaries.

Authorship and reproducibility. This manuscript, code refinements and formalization used AI assistance. Human author review and independent replication are pending. All hypothetical study sizes and mechanism extensions are labeled as proposals; implementation tests are not model trials. No peer review or journal or conference submission is claimed.

References

- [1] Parth Asawa, Christopher M. Glaze, Gabriel Orlanski, Ramya Ramakrishnan, Benji Xu, Asim Biswal, Vincent Sunn Chen, Frederic Sala, Matei Zaharia, and Joseph E. Gonzalez. Continual learning bench: Evaluating frontier AI systems in real-world stateful environments. *arXiv:2606.05661*, 2026. URL <https://arxiv.org/abs/2606.05661>.
- [2] Ali Behrouz, Peilin Zhong, and Vahab Mirrokni. Titans: Learning to memorize at test time. *arXiv:2501.00663*, 2025. URL <https://arxiv.org/abs/2501.00663>.
- [3] Florent Delgrange, Raphael Avalos, and Willem Röpke. Deep SPI: Safe policy improvement via world models. In *Proceedings of ICLR*, 2026. URL <https://arxiv.org/abs/2510.12312>.
- [4] Haoran Ding, Zhaoguo Wang, Yicun Yang, Dexin Zhang, Zhenglin Xu, Haibo Chen, Ruzica Piskac, and Jinyang Li.

- Proving query equivalence using linear integer arithmetic. *Proceedings of the ACM on Management of Data*, 1(4):227:1–227:26, 2023. doi: 10.1145/3626768. URL <https://www.cs.yale.edu/homes/piskac/papers/2024SIGMOD.pdf>.
- [5] Cynthia Dwork, Vitaly Feldman, Moritz Hardt, Toniann Pitassi, Omer Reingold, and Aaron Roth. Generalization in adaptive data analysis and holdout reuse. arXiv:1506.02629, 2015. URL <https://arxiv.org/abs/1506.02629>.
 - [6] Mehrdad Farajtabar, Navid Azizan, Alex Mott, and Ang Li. Orthogonal gradient descent for continual learning. In *Proceedings of AISTATS*, volume 108 of *PMLR*, pages 3762–3773, 2020. URL <https://proceedings.mlr.press/v108/farajtabar20a.html>.
 - [7] Yang He, Pinhan Zhao, Xinyu Wang, and Yuepeng Wang. VeriEQL: Bounded equivalence verification for complex SQL queries with integrity constraints. *Proceedings of the ACM on Programming Languages*, 8(OOPSLA1):132:1–132:29, 2024. doi: 10.1145/3649849. URL <https://arxiv.org/abs/2403.03193>.
 - [8] Steven R. Howard, Aaditya Ramdas, Jon McAuliffe, and Jasjeet Sekhon. Time-uniform, nonparametric, nonasymptotic confidence sequences. *The Annals of Statistics*, 49(2):1055–1080, 2021. doi: 10.1214/20-AOS1991. URL <https://arxiv.org/abs/1810.08240>.
 - [9] Emilie Kaufmann, Olivier Cappé, and Aurélien Garivier. On the complexity of best-arm identification in multi-armed bandit models. *Journal of Machine Learning Research*, 17(1):1–42, 2016. URL <https://jmlr.org/papers/v17/kaufman16a.html>.
 - [10] Romain Laroche, Paul Trichelair, and Rémi Tachet des Combes. Safe policy improvement with baseline bootstrapping. In *Proceedings of ICML*, volume 97 of *PMLR*, 2019. URL <https://proceedings.mlr.press/v97/laroche19a.html>.
 - [11] David Lopez-Paz and Marc’Aurelio Ranzato. Gradient episodic memory for continual learning. arXiv:1706.08840, 2017. URL <https://arxiv.org/abs/1706.08840>.
 - [12] Andrei A. Rusu, Neil C. Rabinowitz, Guillaume Desjardins, Hubert Soyer, James Kirkpatrick, Koray Kavukcuoglu, Razvan Pascanu, and Raia Hadsell. Progressive neural networks. arXiv:1606.04671, 2016. URL <https://arxiv.org/abs/1606.04671>.
 - [13] Yiheng Shu, Bernal Jiménez Gutiérrez, Saisri Padmaja Jonnalagedda, Yuguang Yao, Huan Sun, and Yu Su. AgentCL: Toward rigorous evaluation of continual learning in language agents. arXiv:2606.02461, 2026. URL <https://arxiv.org/abs/2606.02461>.
 - [14] Satinder Singh, Michael James, and Matthew Rudary. Predictive state representations: A new theory for modeling dynamical systems. In *Proceedings of UAI*, pages 512–519, 2004. URL <https://arxiv.org/abs/1207.4167>.
 - [15] Mirac Suzgun, Mert Yuksekgonul, Federico Bianchi, Dan Jurafsky, and James Zou. Dynamic cheatsheet: Test-time learning with adaptive memory. arXiv:2504.07952, 2025. URL <https://arxiv.org/abs/2504.07952>.
 - [16] Arnub Tandon, Karan Dalal, Xinhao Li, Daniel Kocreja, Marcel Rød, Sam Buchanan, Xiaolong Wang, Jure Leskovec, Sanmi Koyejo, Tatsunori Hashimoto, Carlos Guestrin, Jed McCaleb, Yejin Choi, and Yu Sun. End-to-end test-time training for long context. arXiv:2512.23675, 2025. URL <https://arxiv.org/abs/2512.23675>.
 - [17] Philip S. Thomas, Georgios Theodorou, and Mohammad Ghavamzadeh. High confidence policy improvement. In *Proceedings of ICML*, volume 37 of *PMLR*, pages 2380–2388, 2015. URL <https://proceedings.mlr.press/v37/thomas15.html>.
 - [18] Guanzhi Wang, Yuqi Xie, Yunfan Jiang, Ajay Mandlekar, Chaowei Xiao, Yuke Zhu, Linxi Fan, and Anima Anandkumar. Voyager: An open-ended embodied agent with large language models. arXiv:2305.16291, 2023. URL <https://arxiv.org/abs/2305.16291>.
 - [19] Jun Wang. Memento 2: Learning by stateful reflective memory. arXiv:2512.22716, version 2, 2025. URL <https://arxiv.org/html/2512.22716v2>.
 - [20] Zora Zhiruo Wang, Jiayuan Mao, Daniel Fried, and Graham Neubig. Agent workflow memory. arXiv:2409.07429, 2024. URL <https://arxiv.org/abs/2409.07429>.
 - [21] Ian Waudby-Smith and Aaditya Ramdas. Estimating means of bounded random variables by betting. *Journal of the Royal Statistical Society Series B*, 86(1):1–27, 2024. doi: 10.1093/jrsssb/qkad009. URL <https://arxiv.org/abs/2010.09686>.
 - [22] Tianxin Wei, Naveen Sachdeva, Benjamin Coleman, Zhankui He, Yuanchen Bei, Xuying Ning, Mengting Ai, Yunzhe Li, Jingrui He, Ed H. Chi, Chi Wang, Shuo Chen, Fernando Pereira, Wang-Cheng Kang, and Derek Zhiyuan Cheng. Evo-memory: Benchmarking LLM agent test-time learning with self-evolving memory. arXiv:2511.20857, 2025. URL <https://arxiv.org/abs/2511.20857>. Version 2, revised May 2026.
 - [23] Qizheng Zhang, Changran Hu, Shubhangi Upasani, Boyuan Ma, Fenglu Hong, Vamsidhar Kamanuru, Jay Rainton, Chen Wu, Mengmeng Ji, Hanchen Li, Urmish Thakker, James Zou, and Kunle Olukotun. Agentic context engineering: Evolving contexts for self-improving language models. arXiv:2510.04618, 2025. URL <https://arxiv.org/abs/2510.04618>. Version 3, revised March 2026.
 - [24] Huichi Zhou, Siyuan Guo, Anjie Liu, Zhongwei Yu, Ziqin Gong, Bowen Zhao, Zhixun Chen, Menglong Zhang, Yihang Chen, Jinsong Li, Runyu Yang, Qiangbin Liu, Xinlei Yu, Jianmin Zhou, Na Wang, Chunyang Sun, and Jun Wang. Memento-skills: Let agents design agents. arXiv:2603.18743, 2026. URL <https://arxiv.org/abs/2603.18743>.